

## 1) Get the List of Open weight LLM from different service providers

### a) Open ai:

**gpt-oss-120b** (and its safeguard variant): A large-scale model with ~117B parameters (~5.1B active), optimized to fit on a single 80GB GPU like an NVIDIA H100.

**gpt-oss-20b** (and its safeguard variant): A smaller, lower-latency model with ~21B parameters (~3.6B active) built for resource-constrained environments.

### b) Qwen:

**Small & Edge Models:** Highly efficient variants (such as Qwen 3.5 scales down to sub-1B up to mid-sized parameters) engineered to run locally on consumer hardware like laptops and smartphones.

**Qwen2.5 and Qwen3 series: Dense & Mixture-of-Experts (MoE):** Mid-to-large tiers spanning popular configuration sizes like 7B, 14B, 32B, 72B, and massive frontier-class open configurations.

**Qwen3.8-2.4T-A95B:** Alibaba's high-end open-weight release boasting a massive 2.4-trillion total parameter architecture (activating 95B parameters per token) tailored for heavy coding and multi-step agentic tasks.

### c) Deepseek:

**DeepSeek-V4 Pro:** A massive, high-performance Mixture-of-Experts (MoE) model featuring roughly 1.6T total parameters (with ~49B active per token). It is built for heavy reasoning, complex coding, and large-scale agentic workflows.

**DeepSeek-V4 Flash (and V4.1 variants):** A lighter, highly optimized MoE architecture (around 284B–552B total parameters with ~13B–16B active) engineered specifically to drastically cut serving costs and latency while preserving strong multi-step capability

## **2)Hallucination check for the below prompt**

**Write a 4-line poem where every line ends with a word that strictly rhymes with the word 'Silver'.**

### **a)qwen ai**

The golden sun descends beyond the river,  
And makes the weeping willow branches quiver,  
The evening chill will make your fingers shiver,  
As daylight fades into a narrow sliver.

### **b) Chatgpt:**

Strictly speaking, “silver” has no common perfect rhyme in English. So the only strictly exact rhyme is the word itself:

Moonlight whispers softly—silver  
Stars are painted bright—silver  
Dreams drift through the night—silver  
Morning leaves the world—silver

### **c)Deepseek**

No word in English truly rhymes with silver,  
So poets who try are left in a pilver.  
They search and they strain, and they start to shiver,  
Then give up the ghost and go quaff a zilver.

### **d)Gemini**

The moonlit current stirs the stream,  
And sets the water all aquiver,  
It coats the pebbles down below,  
Like polished pieces of a sliver.

3) Model Details with Context window:

Model 1:

 **Inference.net: Schematron V2 Turbo**

inference-net/schematron-v2-turbo

 [Model weights](#)

Compare

Playground

Get API Key

inference-net/schematron-v2-turbo

Schematron V2 Turbo is a 3B-parameter HTML-to-JSON extraction model from Inference.net. It prioritizes throughput for high-volume extraction workloads. Extraction instructions must be supplied through a JSON schema in response\_format rather than through system or user prompts.

MODALITIES  
T → T

IN / OUT PRICE  
\$0.03 / \$0.15 per 1M

CONTEXT  
128K

RELEASED  
Sep 12, 2026

Model 2:

 **Inference.net: Schematron V2 Small**

inference-net/schematron-v2-small

 [Model weights](#)

Compare

Playground

Get API Key

inference-net/schematron-v2-small

Schematron V2 Small is a 3B-parameter HTML-to-JSON extraction model from Inference.net. It prioritizes extraction quality for complex schemas and long pages. Extraction instructions must be supplied through a JSON schema in response\_format rather than through system or user prompts.

MODALITIES  
T → T

IN / OUT PRICE  
\$0.05 / \$0.23 per 1M

CONTEXT  
128K

RELEASED  
Sep 12, 2026

Model 3:

 **OpenAI: GPT-5.6 Luna**

openai/gpt-5.6-luna



Compare

Playground

Get API Key

openai/gpt-5.6-luna

GPT-5.6 Luna is a fast, cost-efficient model in OpenAI's GPT-5.6 series. It is suited for high-volume, latency-sensitive tasks such as chat, classification, and lightweight agentic workflows, providing capable reasoning for its price tier.

MODALITIES  
D G T → T

IN / OUT PRICE  
\$0.20 / \$1.20 per 1M

CONTEXT  
1.1M

RELEASED  
Jul 9, 2026

KNOWLEDGE CUTOFF  
Feb 2026

Model 4:

 **OpenAI: GPT-5.6 Sol**

Compare

Playground

Get API Key

openai/gpt-5.6-sol

GPT-5.6 Sol is the flagship model in OpenAI's GPT-5.6 series. It is suited for complex reasoning, coding, and agentic workflows, and is particularly strong at command-line and multi-step coding tasks and long-horizon problem solving.

|                                             |                                                |                 |                         |                              |
|---------------------------------------------|------------------------------------------------|-----------------|-------------------------|------------------------------|
| MODALITIES<br>Image Video Text Audio → Text | IN / OUT PRICE<br>\$2 / \$10 per 1M<br>50% off | CONTEXT<br>1.1M | RELEASED<br>Jul 9, 2026 | KNOWLEDGE CUTOFF<br>Feb 2026 |
|---------------------------------------------|------------------------------------------------|-----------------|-------------------------|------------------------------|

Model 5:

 **OpenAI: GPT-6 Astra**

Compare

Playground

Get API Key

openai/gpt-6-astra

GPT-6 Astra is OpenAI's flagship model for demanding end-to-end work. It is suited for advanced analysis, software engineering, deep research, scientific work, and document creation, with particular strengths in long-horizon agentic tasks that involve computer and browser use.

|                                             |                                      |                 |                         |
|---------------------------------------------|--------------------------------------|-----------------|-------------------------|
| MODALITIES<br>Image Video Text Audio → Text | IN / OUT PRICE<br>\$10 / \$50 per 1M | CONTEXT<br>1.1M | RELEASED<br>Sep 4, 2026 |
|---------------------------------------------|--------------------------------------|-----------------|-------------------------|

Model 6:

 **OpenAI: gpt-oss-120b**

Compare

Playground

Get API Key

openai/gpt-oss-120b

Model weights

gpt-oss-120b is an open-weight, 117B-parameter Mixture-of-Experts (MoE) language model from OpenAI designed for high-reasoning, agentic, and general-purpose production use cases. It activates 5.1B parameters per forward pass and is optimized to run on a single H100 GPU with native MXFP4 quantization. The mode...  
Show more

|                           |                                          |                 |                         |                              |
|---------------------------|------------------------------------------|-----------------|-------------------------|------------------------------|
| MODALITIES<br>Text → Text | IN / OUT PRICE<br>\$0.03 / \$0.17 per 1M | CONTEXT<br>131K | RELEASED<br>Aug 5, 2025 | KNOWLEDGE CUTOFF<br>Jun 2024 |
|---------------------------|------------------------------------------|-----------------|-------------------------|------------------------------|

Model 7:

 **Google: Gemini 3.8 Flash**

Compare

Playground

Get API Key

google/gemini-3.8-flash

Gemini 3.8 Flash is Google's most intelligent Flash model with significant gains from 3.7 Flash across software engineering, agentic tasks, and multi-step reasoning.

|                                             |                                                     |                 |                         |
|---------------------------------------------|-----------------------------------------------------|-----------------|-------------------------|
| MODALITIES<br>Text Image Video Audio → Text | IN / OUT PRICE<br>\$0.75 / \$3.75 per 1M<br>50% off | CONTEXT<br>1.0M | RELEASED<br>Sep 2, 2026 |
|---------------------------------------------|-----------------------------------------------------|-----------------|-------------------------|

Model 8:

Google: Gemini 3 Flash Preview

google/gemini-3-flash-preview

☆

Gemini 3 Flash Preview is a high speed, high value thinking model designed for agentic workflows, multi turn chat, and coding assistance. It delivers near Pro level reasoning and tool use performance with substantially lower latency than larger Gemini variants, making it well suited for interactive development, long running agent...

Show more ▾

MODALITIES

T G D P C → T

IN / OUT PRICE

\$0.50 / \$3 per 1M

CONTEXT

1.0M

RELEASED

Dec 17, 2025

Compare

Playground

Get API Key

## Model 9:

Google: Gemini 3.7 Flash

google/gemini-3.7-flash

☆

Gemini 3.7 Flash is a multimodal model from Google for fast agentic workflows, coding, and complex multi-step reasoning. It is designed for tasks that require responsive performance and reliable multi-step problem solving.

MODALITIES

T G D P C → T

IN / OUT PRICE

\$0.75 / \$3.75 per 1M 

50% off

CONTEXT

1.0M

RELEASED

Aug 13, 2026

Compare

Playground

Get API Key

## Model 10:

Google: Gemini 2.5 Flash Lite

google/gemini-2.5-flash-lite

☆

Gemini 2.5 Flash-Lite is a lightweight reasoning model in the Gemini 2.5 family, optimized for ultra-low latency and cost efficiency. It offers improved throughput, faster token generation, and better performance across common benchmarks compared to earlier Flash models. By default, "thinking" (i.e. multi-pass reasoning) is disabled to prioritize speed, but developers can enable it via the [Reasoning API parameter](#) to selectively trade off cost for intelligence.

Show less ^

MODALITIES

T G D P C → T

IN / OUT PRICE

\$0.10 / \$0.40 per 1M

CONTEXT

1.0M

RELEASED

Jul 22, 2025

KNOWLEDGE CUTOFF

Jan 2025

Compare

Playground

Get API Key